
2018



Q.) How to process 1 TB of data in Spark?

Here this 1TB data is stored in HDFS,
in HDFS the Block size is = 128 MB

$$1\text{TB} = (1 \times 1024 \times 1024) / 128 = 8192 \text{ blocks}$$

Let us take

→ 20 Executors

→ 5 cores in each executor

→ 6 GB RAM in each executor

And the cluster can perform ($20 \times 5 = 100$) parallel tasks. Here tasks mean block so 100 blocks can be processed parallelly.

$$100 \times 128 \text{ MB} = 12800 \text{ MB} / 1024 = 12.5 \text{ GB}$$

(12.5 GB data will get processed in 1st batch)

Since the RAM size is 6 GB in each executor

$$(20 \text{ Exec} \times 6 \text{ GB} = 120 \text{ GB Total RAM})$$

So at a time 12.5 GB of RAM will be occupied in the cluster.

$20 \text{ Node} / 12 \text{ GB} = 1.6 \text{ GB RAM}$ from each executor will be in use.

Now, Available RAM in each executor ($6 \text{ GB} - 1.6 \text{ GB} = 4.4 \text{ GB}$) will be free to use by other user's job.

So,
$$\begin{aligned} LTB &= 1024 \text{ GB} / 12 \text{ GB} \\ &= 85 \text{ Batches} \end{aligned}$$

Whole data will get processed in around 85 Batches



Failure Scenarios in Spark

① Out of Memory Error

- Increase executor memory
- Check if broadcast variable is Big in Size
- Avoid unnecessary data shuffle
- If driver is out of memory then check

if collect() action was called

- ② Network Error: It happens when Spark driver & executor are not able to communicate with each other.

Solution { spark.network.timeout
 spark.executor.heartbeatInterval

- ③ Data Input/Output

↳ Make sure data is consistent as per data contract
↳ Always put read/write related operations in try-except block for easy debugging.

- ④ Serialization Error: Data is not according to the serialized schema.

- ⑤ Disk Space Error: Spark job spills intermediate data into the local

machine disk. If no sufficient memory is there then job will fail.

Solution: spark.local.dir

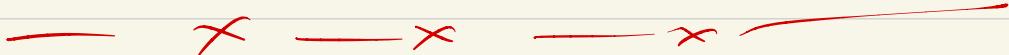
⑥ Memory Configuration

If we are facing heap space related issues and memory overhead issues then try to configure memory components effectively.

→ Analyze the data volume.

→ Define driver & executor memory properly.

→ Adjust memory configuration params properly.



Frequently used optimizations in Spark

① Partitioning:

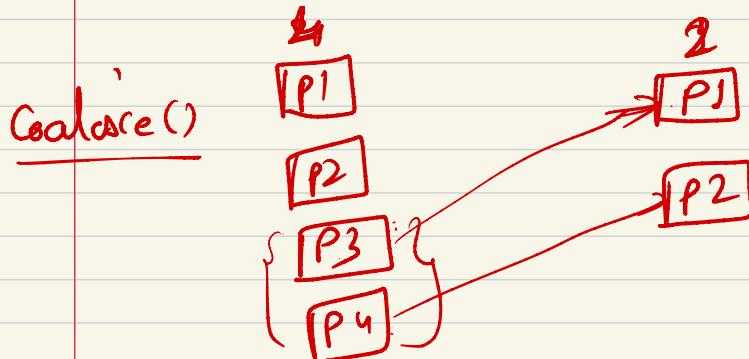
↳ repartition() ↳ Increase
↳ coalesce() ↳ Decrease

`empDF = spark.read.csv("employee.csv")`

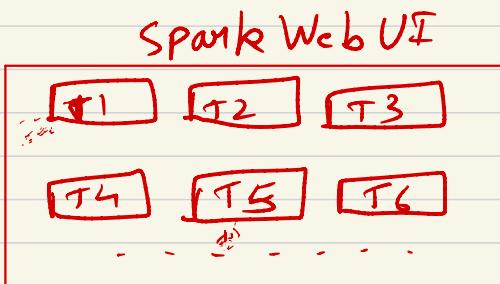
We are reading from HDFS, so let say
20 Block = 20 partitions.

In the code itself, we want to
increase the number of partitions
We can use `repartition()`.

`empDF.repartition(40)`



~~`empDF.coalesce(1).write()`~~



\Rightarrow 60 partitions
 \Rightarrow 60 parallel tasks

`joinDF.repartition(100)`

② Optimization in case of skewness

We can use Key-Salting technique.

③ Broadcast Variables: It will help to broadcast dataset on multiple executors.

④ Persist() and Cache():

- ↳ Cache() → only in memory save
- ↳ Persist() → Memory + Disk both

⑤ Try to minimize number of stages in the spark job. That means unoptimized use of wide transformation.

* Difference between Job vs Stage vs Task?

Job → As soon as we call an action
In Spark application, new Job will be created.

Stage → Wherever data shuffling takes place in the spark app, a new stage will get created

Task → Small unit of execution.

empDF → ①

deptDF → ②

empDF.map()] → Stage 0

empDF.filter('INDIA')] → Stage 1

empDF.groupBy('dept_id').sum('Salary')] → Stage 1

empDF.filter("Salary_Sum < 50000")]

{ empDF.collect(5) }

joinPF = empDF.join(deptDF)

{ joinDF.filter(deptDF.dept_id not null) }

{ joinDF.collect(10) }

Trans

Action

Trans

Action

